🎬 VIDEO DEMO SCRIPT

Sniffer
Live Dashboard Demo

A walkthrough of the Hallucination Detection live web dashboard — all five tabs, explained for camera.

TOTAL RUNTIME	SCENES	DASHBOARD TABS	TONE
~4 minutes	6	5	Demo / Research

Video Timeline Overview

0:00	0:30	1:30	2:20	3:00	3:30	3:50
Intro	Playground	Leaderboard	Experiments	Taxonomy	Literature	Outro

01 Introduction — What is Sniffer?

0:00 – 0:30

SCREEN

Dashboard open in browser. Sidebar visible on the left. "Live Hallucination Playground" is the active tab. Do not interact yet — just let the UI sit.

YOU SAY

> *"Hey everyone — this is **Sniffer**, a live research dashboard I built for my empirical NLP study on hallucination detection."*
>
> *"Large language models are powerful, but they often generate text that sounds confident yet is completely wrong. Sniffer lets you test LLM-generated responses in real time against four different detection strategies, and compare how well each one performs."*
>
> *"Let me walk you through it."*

02 Tab 1 — Live Hallucination Playground

0:30 – 1:30

SCREEN

The Playground tab is already active. Pan slowly across the three input fields on the left panel.

YOU SAY

> *"The first tab is the Live Playground. You give it three things: the original **prompt or question**, the **grounding context** — think of this as the 'truth' document — and then the **LLM-generated response** you want to check."*

ACTION

Click the `Contradiction` preset chip to auto-fill a sample where the LLM gets the founding year wrong. Read the inputs out briefly.

YOU SAY

> *"I'll load this pre-built sample — the context says Acme was founded in **1998**, but the LLM response confidently says **2005**. That's a classic factual contradiction hallucination. Let me hit Detect."*

ACTION

Click the **Detect Hallucinations** button. Wait for results. Point to the right panel as it populates.

YOU SAY

> *"On the right you get two things. First, a **FActScore-style atomic claim decomposition** — the response is broken into individual verifiable claims, each tagged as Supported, Contradicted, or Unsupported against the source."*
>
> *"Below that is the **Detector Comparison Grid** — four cards, one for each detection strategy. Each shows its hallucination confidence score, its latency in milliseconds, and its estimated cost per thousand queries. You can immediately see which detectors agree and which disagree."*

PRO TIP

Try the `Factual Baseline` and `Unsupported Info` chips too if you have time — great for showing contrast.

03 Tab 2 — Detector Leaderboard

1:30 – 2:20

ACTION

Click **Detector Leaderboard** in the left sidebar. Let the table render fully.

YOU SAY

> *"Now let's look at the Leaderboard. This is the empirical evaluation — all four detectors run against real benchmark samples drawn from **HaluEval, RAGTruth, SummaC, and TruthfulQA**."*
>
> *"The table gives us Accuracy, Precision, Recall, F1 Score, AUROC, and Latency side by side. The highlighted row is the **Ensemble Meta-Classifier**, which combines signals from the other three detectors."*
>
> *"What's interesting here is the trade-off — **SelfCheckGPT** has strong recall but it's slow and expensive, whereas the **Similarity Baseline** is lightning fast but misses nuanced unsupported claims. The Ensemble tries to capture the best of all worlds."*

ACTION

Slowly scroll to the **Cost / 1K Queries** column and let it sit on screen while you finish the quote.

04 Tab 3 — Research Experiments

2:20 – 3:00

ACTION

Click **Research Experiments**. Two charts will appear — the Pareto Frontier scatter plot and the RAG Ablation bar chart.

YOU SAY

> *"This tab houses the deeper empirical experiments."*
>
> *"The first chart is a **Pareto Frontier** — Cost on the x-axis in log scale, F1 score on the y-axis. Each dot is a detector. An ideal detector sits in the top-left — highest accuracy at lowest cost. This chart makes it visually obvious whether spending a hundred times more compute is actually worth it."*
>
> *"The second chart is a **RAG Ablation study**. It tests all four detectors under three conditions: no context, partial context, and full retrieval-augmented context. As you'd expect, all detectors improve dramatically with better grounding — but not all of them by the same amount."*

PRO TIP

Hover over the Pareto scatter dots to show tooltips if they're visible — it makes the chart feel more interactive on camera.

05 Tab 4 — Failure Taxonomy Matrix

3:00 – 3:30

ACTION

Click **Failure Taxonomy**. Scroll down the category matrix table slowly so every row is visible for a moment.

YOU SAY

> *"The Failure Taxonomy tab is one of my favourites. It breaks hallucination types into **six fine-grained categories** — things like factual contradictions, unsupported claims, temporal errors, entity confusions, and so on."*
>
> *"For each category, the matrix shows the accuracy of every detector. This is where the nuance really shows — the NLI detector is great at catching direct contradictions, but really struggles with unsupported-but-technically-true claims. The Ensemble handles both significantly better."*

06 Tab 5 — Literature & Outro

3:30 – 4:00

ACTION

Click **Core Literature (30 Papers)**. Scroll slowly through the paper cards. Briefly hover on a starred paper to highlight it.

YOU SAY

> *"Finally, the Literature tab curates the **30 core papers** this study is grounded in — spanning ACL, EMNLP, NAACL and Nature publications. The starred*

> *papers are the most foundational references, and clicking through takes you to the original ACL Anthology pages."*
>
> *"This traces the field from early factual consistency metrics all the way to modern zero-resource detection approaches — giving the whole project proper academic grounding."*
>
> *"That's the Sniffer dashboard — a fully interactive research tool for understanding and evaluating hallucination detection in LLMs. Thanks for watching!"*